

Design and Analysis of Algorithms

L23: Job Scheduling

Dr. Ram P Rustagi
Sem IV (2019-H1)
Dept of CSE, KSIT/KSSEM
rprustagi@ksit.edu.in

Resources

- Text book 2: Sec 4.1, 4.3, 4.4
- Text book 1: Sec 9.1-5.4 - Levitin
- RI: Introduction to Algorithms
 - Cormen et al.

Example Case

- In college fest which starts at 9:00am, there are a number of available events as below to participate, and each event takes 1 unit of time (e.g. 1hr).
 - Each event has different awards values
 - Each event has its own closing timeline.

| Event | Closing | Award |
|----------|---------|-------|
| Mimcry | 12:00 | 200 |
| Drama | 11:00 | 100 |
| Painting | 12:00 | 90 |
| Dance | 10:00 | 50 |
| JAM | 11:00 | 125 |
| Singing | 10:00 | 60 |

| Deadline |
|----------|
| 3 |
| 2 |
| 3 |
| 1 |
| 2 |
| 1 |

Q: What is the max award you can get?

Greedy Job Scheduling

- A set of n jobs to run on a computer
- Each job i has a deadline $d_i \geq 1$ and profit $p_i \geq 0$
- There is only one computer
- Each job takes one unit of time (simplification)
- Profit is earned when job is completed by deadline
- Find the subset of jobs that maximizes the profit, i.e.

$$\text{Maximize } \sum_{i \in J} P_i$$

Note: It belongs to subset paradigm since we are looking at subset of jobs.

Example: Job Scheduling

| Job | Profit | Dead-line |
|-----|--------|-----------|
| 1 | 100 | 2 |
| 2 | 10 | 1 |
| 3 | 15 | 2 |
| 4 | 27 | 1 |

Optimal Solution: 1,4

| Feasible Solutions | Profit |
|--------------------|--------|
| 1 | 100 |
| 2 | 10 |
| 3 | 15 |
| 4 | 27 |
| 1,2 | 110 |
| 1,3 | 115 |
| 1,4 | 127 |
| 2,3 | 25 |
| 3,4 | 42 |

Job Scheduling: Greedy Approach

- What should be the optimization measure to schedule the next job?
- First attempt:
 - Choose $\sum_{i \in J} P_i$ as the optimization measure
 - i.e. choose a job that increases this value maximum
 - Subject to constraint of the deadline i.e. J (set of jobs) should be feasible solution.
 - How to choose jobs:
 - Order jobs in decreasing order of profit
 - Choose job one at a time as per this order and add to the solution if solution remains feasible.

Job Scheduling: Greedy Approach

| Job | Pro fit | De ad- line |
|-----|------------|-------------------|
| 1 | 100 | 2 |
| 2 | 10 | 1 |
| 3 | 15 | 2 |
| 4 | 27 | 1 |

- Application of First Greedy approach
 - Job 1 is added to J. Feasible $\{1\}$
 - Next: Job 4 is considered as per order.
 - Is set $J = \{1, 4\}$ feasible.
 - Yes if 4-1, No if 1-4
 - Thus $\{1, 4\}$ is feasible solution.
 - Next: Job 3 is considered,
 $\{1, 4, 3\}$ is infeasible, thus J remains $\{1, 4\}$
 - Next: Job 2 is considered
 $\{1, 4, 2\}$ is infeasible thus J remains $\{1, 4\}$
 - The max profit is 127 for $J = \{1, 4\}$
- Time complexity :
 - $n!$ to evaluate feasibility for a given set

Job Scheduling: Feasible Solution

- How to determine that a given set of jobs constitute feasible solution.
- Try out all possible permutations in jobs J
 - Check for each permutation if jobs can be scheduled meeting the deadlines.
- Easy to check or a given permutation $\sigma = i_1, i_2, \dots, i_k$
 - Job i_q must be completed by time q , $1 \leq q \leq k$
 - If for some job i_q , $q > d_{i_q}$, then job i_q is not completed by d_{i_q} .
- When $|J| = k$, all $k!$ permutations must be checked
- Can we find one permutation that meets the need?
 - Order the jobs in non-decreasing order of deadlines

Proof for Feasible Solution

- Theorem 1:
 - Let J be the set of k jobs and $\sigma = i_1, i_2, \dots, i_k$ is a permutation of jobs in J such that $d_{i_1} \leq d_{i_2} \leq \dots \leq d_{i_k}$.
Then J is a feasible solution if and only if (*iff*) the jobs in J can be processed in the order σ without violating any deadline.
- Theorem 2:
 - The greedy method describes above always obtains an optimal solution to the job scheduling problem.

Algo High Level

```
Algo GreedyJob(int d[], set J, int n) {  
    // J is set of jobs that can be completed in deadlines d[]  
    J = {1}  
    for i = 2 to n {  
        if all jobs in J ∪ {i} can be completed, then  
            // by their deadlines  
            J = J ∪ {i}  
    }  
}
```

Algo-I: Job Scheduling

```
int JobSchedule2(int d[], int j[], int n) {  
    //  $n \geq 1$ , and deadlines  $d[i] \geq 1, 1 \leq i \leq n$   
    // Jobs are ordered such that their profits are in non-  
    // increasing order i.e.  $p[1] \geq p[2] \geq \dots \geq p[n]$ .  
    //  $J[i]$  is the  $i^{\text{th}}$  job in the optimal solution with  $k \leq n$  jobs  
    // At algo termination,  $d[J[i]] \leq d[j[i+1]]$ ,  $1 \leq i < k$   
  
    // Initialize  
    d[0] = 0 // fictitious job with deadline of 0  
    // allows for job insertion at position 1 later.  
    J[0] = 0 // this job is boundary and can't be scheduled  
    J[1] = 1 // start with job 1 with highest profit  
    k = 1 // job set size is 1 to start with
```

Algo1: Job Scheduling

```
for (i=2; i≤n; i++) {  
    // consider jobs in non-increasing order of p[i]  
    // find pos for J[i] and check for feasibility of insertion  
    int r = k //job set size  
    while ( (d[J[r]]>d[i]) && (d[J[r]]!=r) )  
        r--; //find position where job i can be considered.  
    if ( (d[J[r]]≤d[i]) && (d[J[r]]>r) ) {  
        //insert i into J[]  
        for (int q=k; q≥(r+1); q--)  
            J[q+1]=J[q] // increase deadline of jobs by 1.  
        J[r+1]=i  
        k++ // since job i is feasible, increase the set size.  
    } //end if  
} //end for i  
return k  
}
```

Algo-1: Time Complexity

- For loop run n times.
 - Each job needs to be considered.
- if K is the value of max deadline, then
 - Inside while loop plus for loop (for shifting slots) may run K times.
- Time complexity: $O(nK)$
- Considering K is of order of n (if all jobs can be scheduled)
- Time complexity: $O(n^2)$

Algo-2: Job Scheduling

```
//Approach: schedule a job in the slot where it meets deadline.  
// If no slot is available before deadline, then job is not scheduled.  
// jobs are ordered in non-increasing order as per deadlines.  
int JobSchedule-1(int d[], int j[], int n) {  
    //  $n \geq 1$ , and deadlines  $d[i] \geq 1, 1 \leq i \leq n$   
    // Jobs are ordered such that their profits are in non-  
    // increasing order i.e.  $p[1] \geq p[2] \geq \dots \geq p[n]$  .  
    // Job[i] is  $i^{\text{th}}$  job in the optimal solution with  $k \leq n$  jobs  
    // At algo termination,  $d[\text{Job}[i]] \leq d[\text{job}[i+1]]$  ,  
    //  $1 \leq i < k$   
    // Initialization  
     $k = 0$  ; // size of Job schedule  
    for  $i = 1$  to  $n$   
        slot[i] = False // all slots are initialized to false
```

Algo-2: Job Scheduling

```
for (i=1; i≤n; i++) {  
    // consider jobs in non-increasing order of p[i]  
    //check if any slot available before deadline  
    while (j=d[i]; j>0; j-) {  
        //find position where job i can be considered.  
        if (slot[j] == False{  
            //Add jobs to the slot  
            slot[j] = True;  
            Job[j] = i;  
            k++;  
            break  
        } //end if  
    } //end while  
} // end for  
return k  
}
```

Algo-2: Time Complexity

- For loop run n times.
 - Each job needs to be considered.
- if K is the value of max deadline, then
 - while loop may run K times.
- Time complexity: $O(nK)$
- Considering K is of order of n (if all jobs can be scheduled)
- Time complexity: $O(n^2)$

Union-Find Approach

- Using set based Union-Find approach
 - It is almost $O(n)$
- Union-Find approach
- Consider the set of n elements which are known.
- All these elements put in an array and their id can be the array index i.e.
 - $a[i] = x_i$ # i^{th} element is s_i .
- Elements are divided into groups (sets)
 - Initially, each element is a group by itself.
- Two kinds of operations:
 - Find the group to which element belongs
 - Merge the two groups

Union-Find Approach

- Two operations given below are performed in arbitrary order
 - *Find(i)*: return the group id containing element x_i .
 - *Union(A, B)*: Combine the (set) group A with (set) group B to form a new group.
 - Give a unique name to this group. All elements of this new group will have this group id.
 - This could be one of earlier groups as well i.e.
 - The new names should conflict with other names.
- Goal: Design an efficient data structure that will support any sequence of these two operations.

Union-Find Approach

- Approach 1:
 - Keep Find(i) efficient. Since all elements are accessible at the i^{th} index in array,
 - This can take $O(1)$ time. Essentially, a trivial operation.
 - Union(A, B) is expected to take more time.
 - Either change the id of all elements of A to that of B or vice versa.
 - Typically, take the smaller set and change group identity of these elements to that of larger set.
- Time complexity: $O(n \log_2 n)$
 - Each time an element's group is changed, group size at least doubles.

Union-Find Approach: Improved

- Approach 2:
 - Make Union(A, B) efficient at the cost of Find(i).
 - Make Union operation takes constant time and improve upon the time taken by Find.
 - Use the indirect addressing for union to make it in $O(1)$ time.
 - Each entry in the array has two parts
 - Identity of element i.e. group id or value.
 - Pointer which is initially `Nil`.
 - Union(A, B) is performed by making the pointer of B to point to A .
 - After several union operations, data structures becomes like a tree.

Union-Find Approach: Improved





Fig: Representation for Union Find problem

- The element at root of the Tree is the identify of group.
- To find the group of an element, follow the path till the root of the tree.

Union-Find Approach: Improved



- Take the tree with smaller number of nodes to point to root of the tree of larger number of nodes.
- This requires storing the number of elements in the tree at the root as well.
- On Union operation, update the pointer of smaller tree and count at the larger tree.
 - Break the tie arbitrarily

Union-Find Approach: Improved

- Basic idea: Balance and collapse the tree
 - Union operation still takes $O(1)$ time
 - Changing the pointer $O(1)$ time
 - Updating the count $O(1)$ time
 - Thus, $O(1) + O(1) = O(1)$

Union-Find Approach: Improved

- Theorem: When balancing is used, the tree of height h will contain at least 2^h nodes.
- Proof outline
 - First union operation results in tree of height 1 with two elements.
 - Consider A is height h_A and B is of height h_B .
 - Let A is larger tree. Thus, on merging B, root of B points to root of A.
 - If $h_A > h_B$, then A's height remains h_A i.e. unchanged
 - Otherwise, height of tree becomes $h_B + 1$
 - Thus, with increase in height of 1, the size of tree has at least doubled.
- Thus, time taken for a Find(X) operation is $O(\log_2 n)$

Union-Find Approach: Improved

- Further improvement
- Any time we do a find operation, change the pointers of all the nodes in the path to directly point to the root of the tree.
 - This is called **path compression**.
- Traversing the path takes only double the number of steps, and thus
 - Time complexity of find remains the same.
- Time complexity with path compression for m operations is given by
 $O(m \log^* n)$, where $\log^* n$ is iterated logarithm function

Iterated Logarithm

- Iterated logarithm function is defined as

$$\log^* n = 1 + \log^* (\lceil \log_2 n \rceil)$$

$$\log^* 2 = 1 \quad (\text{Given})$$

$$\log^* 4 = \log^* 2^2$$

$$= 1 + \log^* (\lceil \log_2 2^2 \rceil)$$

$$= 1 + \log^* 2 = 1 + 1 = 2$$

$$\log^* 16 = \log^* 2^4$$

$$= 1 + \log^* (\lceil \log_2 2^4 \rceil)$$

$$= 1 + \log^* 4 = 1 + 2 = 3$$

$$\log^* 65536 = \log^* 2^{16}$$

$$= 1 + \log^* (\lceil \log_2 2^{16} \rceil)$$

$$= 1 + \log^* 16 = 1 + 4 = 5$$

$$\log^* 2^{65536} = 1 + \log^* 65536 = 5 \quad (\text{very large } n)$$

Fast Job Scheduling (Union-Find)

- Let i denote the timeslot i
 - At the start time, each time slot is its own set
- There are m timeslots, where
$$m = \min(n, \max(d_i))$$
 - i.e. the latest deadline
- Each set of k slots has a value $F(k)$ for all slots i in set k
 - $F(k)$: Stores highest free timeslot before this time
 - $F(k)$: Defined only for root node in set
- Initially all slots are free

Summary

- ?